

---

FEBRUARY 27, 2012 | BAPTISTE ASSMANN | 41 COMMENTS

We've seen recently more and more DOS and DDOS attacks. Some of them were very big, requiring thousands of computers...

But in most cases, this kind of attacks are made by a few computers aiming to make a service or website unavailable, either by sending it too many requests or by taking all its available resources, preventing regular users to use the service.

Some attacks targets known vulnerabilities of widely used applications.

In the present article, we'll explain how to take advantage of an **application delivery controller** to protect your website and application against DOS, DDOS and vulnerability scans.

Why using a LB for such protection since a firewall and a Web Application Firewall (aka WAF) could already do the job?

Well, the Firewall is not aware of the application layer but would be useful to protect against SYN flood attacks. That's why we saw recently application layer firewalls: Web Application Firewalls, also known as WAF.

Well, since the load balancer is in front of the platform, it can be a good partner for the WAF, filtering out 99% of the attacks, which are managed by script kiddies. The WAF can then happily clean up the remaining attacks.

Well, maybe you don't need a WAF and you want to take advantage of your **Aloha** and save some money ;).

Note that **you need an application layer load-balancer**, like **Aloha** or OpenSource **HAProxy** to be efficient.

## TCP syn flood attacks

The syn flood attacks consist in sending as many TCP syn packets as possible to a single server trying to saturate it or at least, saturating its uplink bandwidth.

---

```
1 # Protection SYN flood
2 net.ipv4.tcp_syncookies = 1
3 net.ipv4.conf.all.rp_filter = 1
4 net.ipv4.tcp_max_syn_backlog = 1024
```

---

**Note:** If the attack is very big and saturates your internet bandwidth, the only solution is to ask your internet access provider to null route the attackers IPs on its core network.

## Slowloris like attacks

For this kind of attack, the clients will **send very slowly their requests to a server**: header by header, or even worst character by character, waiting long time between each of them.

The server have to wait until the end of the request to process it and send back its response.

The purpose of this attack is to prevent regular users to use the service, since the attacker would be using all the available resources with very slow queries.

In order to protect your website against this kind of attack, just setup the **HAProxy** option “**timeout http-request**”.

You can set it up to 5s, which is long enough.

It tells **HAProxy** to let five seconds to a client to send its whole HTTP request, otherwise **HAProxy** would shut the connection with an error.

For example:

---

```
1 # On Aloha, the global section is already setup for you
2 # and the haproxy stats socket is available at /var/run/haproxy.stats
3 global
4     stats socket ./haproxy.stats level admin
5
6 defaults
7     option http-server-close
8     mode http
9     timeout http-request 5s
```

---

```
18 stats uri /
19 stats realm HAProxy Statistics
20 stats auth admin:admin
21
22 frontend ft_web
23 bind 0.0.0.0:8080
24
25 # Spalreadylit static and dynamic traffic since these requests have diffe
26 use_backend bk_web_static if { path_end .jpg .png .gif .css .js }
27
28 default_backend bk_web
29
30 # Dynamic part of the application
31 backend bk_web
32 balance roundrobin
33 cookie MYSRV insert indirect nocache
34 server srv1 192.168.1.2:80 check cookie srv1 maxconn 100
35 server srv2 192.168.1.3:80 check cookie srv2 maxconn 100
36
37 # Static objects
38 backend bk_web_static
39 balance roundrobin
40 server srv1 192.168.1.2:80 check maxconn 1000
41 server srv2 192.168.1.3:80 check maxconn 1000
```

To test this configuration, simply open a **telnet** to the frontend port and wait for 5 seconds:

```
telnet 127.0.0.1 8080
Trying 127.0.0.1...
Connected to 127.0.0.1.
Escape character is '^]'.
HTTP/1.0 408 Request Time-out
Cache-Control: no-cache
Connection: close
Content-Type: text/html

<h1>408 Request Time-out</h1>
Your browser didn't send a complete request in time.

Connection closed by foreign host.
```

- 
- too many connections opened
  - new connection rate too high
  - http request rate too high
  - bandwidth usage too high
  - client not respecting RFCs (IE for SMTP)

## How does a regular browser works?

Before trying to protect your website from weird behavior, we have to define what a “normal” behavior is!

This paragraphs gives the main lines of how a browser works and there may be some differences between browsers.

So, when one wants to browse a website, we use a browser: Chrome, Firefox, Internet Explorer, Opera are the most famous ones.

After typing the website name in the URL bar, the browser will look like for the IP address of your website.

Then it will establish a tcp connection to the server, downloading the main page, analyze its content and follow its links from the HTML code to get the objects required to build the page: javascript, css, images, etc...

To get the objects, **it may open up to 6 or 7 TCP connections per domain name.**

Once it has finished to download the objects, it starts aggregating everything then print out the page.

## Limiting the number of connections per users

As seen before, a browser opens up 5 to 7 TCP connections to a website when it wants to download objetcs and they are opened quite quickly.

One can consider that somebody having more than 10 connections opened is not a regular user.

The configuration below shows how to do this limitation in the **Aloha** and **HAProxy**:

**This configuration also applies to any kind of TCP based application.**

```
5
6 defaults
7     option http-server-close
8     mode http
9     timeout http-request 5s
10    timeout connect 5s
11    timeout server 10s
12    timeout client 30s
13
14 listen stats
15     bind 0.0.0.0:8880
16     stats enable
17     stats hide-version
18     stats uri /
19     stats realm HAProxy Statistics
20     stats auth admin:admin
21
22 frontend ft_web
23     bind 0.0.0.0:8080
24
25     # Table definition
26     stick-table type ip size 100k expire 30s store conn_cur
27
28     # Allow clean known IPs to bypass the filter
29     tcp-request connection accept if { src -f /etc/haproxy/whitelist.lst }
30     # Shut the new connection as long as the client has already 10 opened
31     tcp-request connection reject if { src_conn_cur ge 10 }
32     tcp-request connection track-sc1 src
33
34     # Split static and dynamic traffic since these requests have different im
35     use_backend bk_web_static if { path_end .jpg .png .gif .css .js }
36
37     default_backend bk_web
38
39 # Dynamic part of the application
40 backend bk_web
41     balance roundrobin
42     cookie MYSRV insert indirect nocache
43     server srv1 192.168.1.2:80 check cookie srv1 maxconn 100
44     server srv2 192.168.1.3:80 check cookie srv2 maxconn 100
45
46 # Static objects
47 backend bk_web_static
48     balance roundrobin
49     server srv1 192.168.1.2:80 check maxconn 1000
50     server srv2 192.168.1.3:80 check maxconn 1000
```

---

## Testing the configuration

run an apache bench to open 10 connections and doing request on these connections:

```
ab -n 50000000 -c 10 http://127.0.0.1:8080/
```

Watch the table content on the haproxy stats socket:

```
echo "show table ft_web" | socat unix:./haproxy.stats -  
# table: ft_web, type: ip, size:102400, used:1  
0x7afa34: key=127.0.0.1 use=10 exp=29994 conn_cur=10
```

Let's try to open an eleventh connection using telnet:

```
telnet 127.0.0.1 8080  
Trying 127.0.0.1...  
Connected to 127.0.0.1.  
Escape character is '^]'.  
Connection closed by foreign host.
```

Basically, opened connections can keep on working, while a new one can't be established.

## Limiting the connection rate per user

In the previous chapter, we've seen how to protect ourselves from somebody who wants to open more than X connections at the same time.

Well, this is good, but something which may kill performance would be to allow somebody to open and close a lot of tcp connections over a short period of time.

As we've seen previously, a browser will open up to 7 TCP connections in a very short period of time (a few seconds). One can consider that somebody having more than 20 connections opened over a

```
1 # On Aloha, the global section is already setup for you
2 # and the haproxy stats socket is available at /var/run/haproxy.stats
3 global
4     stats socket ./haproxy.stats level admin
5
6 defaults
7     option http-server-close
8     mode http
9     timeout http-request 5s
10    timeout connect 5s
11    timeout server 10s
12    timeout client 30s
13
14 listen stats
15     bind 0.0.0.0:8880
16     stats enable
17     stats hide-version
18     stats uri /
19     stats realm HAProxy Statistics
20     stats auth admin:admin
21
22 frontend ft_web
23     bind 0.0.0.0:8080
24
25     # Table definition
26     stick-table type ip size 100k expire 30s store conn_rate(3s)
27
28     # Allow clean known IPs to bypass the filter
29     tcp-request connection accept if { src -f /etc/haproxy/whitelist.lst }
30     # Shut the new connection as long as the client has already 10 opened
31     tcp-request connection reject if { src_conn_rate ge 10 }
32     tcp-request connection track-sc1 src
33
34     # Split static and dynamic traffic since these requests have different im
35     use_backend bk_web_static if { path_end .jpg .png .gif .css .js }
36
37     default_backend bk_web
38
39 # Dynamic part of the application
40 backend bk_web
41     balance roundrobin
42     cookie MYSRV insert indirect nocache
43     server srv1 192.168.1.2:80 check cookie srv1 maxconn 100
44     server srv2 192.168.1.3:80 check cookie srv2 maxconn 100
45
46 # Static objects
```

---

have a negative impact for them. You can whitelist these IPs.

## Testing the configuration

run 10 requests with ApacheBench, everything may be fine:

```
ab -n 10 -c 1 -r http://127.0.0.1:8080/
```

Using socat we can watch this traffic in the stick-table:

```
# table: ft_web, type: ip, size:102400, used:1  
0x11faa3c: key=127.0.0.1 use=0 exp=28395 conn_rate(3000)=10
```

Running a telnet to run a eleventh request and the connections get closed:

```
telnet 127.0.0.1 8080  
Trying 127.0.0.1...  
Connected to 127.0.0.1.  
Escape character is '^]'.  
Connection closed by foreign host.
```

## Limiting the HTTP request rate

Even if in the previous examples, we were using HTTP as the protocol, we based our protection on layer 4 information: number or opening rate of TCP connections.

An attacker could respect the number of connection we would set by emulating the behavior of a regular browser.

Now, let's go deeper and see what we can do on HTTP protocol.



```
4 stats socket ./haproxy.stats level admin
5
6 defaults
7   option http-server-close
8   mode http
9   timeout http-request 5s
10  timeout connect 5s
11  timeout server 10s
12  timeout client 30s
13
14 listen stats
15   bind 0.0.0.0:8880
16   stats enable
17   stats hide-version
18   stats uri /
19   stats realm HAProxy Statistics
20   stats auth admin:admin
21
22 frontend ft_web
23   bind 0.0.0.0:8080
24
25   # Use General Purpose Counter (gpc) 0 in SC1 as a global abuse counter
26   # Monitors the number of request sent by an IP over a period of 10 second
27   stick-table type ip size 1m expire 10s store gpc0,http_req_rate(10s)
28   tcp-request connection track-sc1 src
29   tcp-request connection reject if { src_get_gpc0 gt 0 }
30
31   # Split static and dynamic traffic since these requests have different im
32   use_backend bk_web_static if { path_end .jpg .png .gif .css .js }
33
34   default_backend bk_web
35
36   # Dynamic part of the application
37   backend bk_web
38     balance roundrobin
39     cookie MYSRV insert indirect nocache
40
41   # If the source IP sent 10 or more http request over the defined period,
42   # flag the IP as abuser on the frontend
43   acl abuse src_http_req_rate(ft_web) ge 10
44   acl flag_abuser src_inc_gpc0(ft_web)
45   tcp-request content reject if abuse flag_abuser
46
47   server srv1 192.168.1.2:80 check cookie srv1 maxconn 100
48   server srv2 192.168.1.3:80 check cookie srv2 maxconn 100
49
50   # Static objects
51   backend bk_web_static
52     balance roundrobin
```

---

## Testing the configuration

run 10 requests with ApacheBench, everything may be fine:

```
ab -n 10 -c 1 -r http://127.0.0.1:8080/
```

Using socat we can watch this traffic in the stick-table:

```
# table: ft_web, type: ip, size:1048576, used:1  
0xbebbb0: key=127.0.0.1 use=0 exp=8169 gpc0=1 http_req_rate(10000)=10
```

Running a telnet to run a eleventh request and the connections get closed:

```
telnet 127.0.0.1 8080  
Trying 127.0.0.1...  
Connected to 127.0.0.1.  
Escape character is '^]'.  
Connection closed by foreign host.
```

## Detecting vulnerability scans

Vulnerability scans could generate different kind of errors which can be tracked by **Aloha** and **HAProxy**:

- invalid and truncated requests
- denied or tarpitted requests
- failed authentications
- 4xx error pages

HAProxy is able to monitor an error rate per user then can take decision based on it.

```
8 mode http
9 timeout http-request 5s
10 timeout connect 5s
11 timeout server 10s
12 timeout client 30s
13
14 listen stats
15 bind 0.0.0.0:8880
16 stats enable
17 stats hide-version
18 stats uri /
19 stats realm HAProxy Statistics
20 stats auth admin:admin
21
22 frontend ft_web
23 bind 0.0.0.0:8080
24
25 # Use General Purpose Counter 0 in SC1 as a global abuse counter
26 # Monitors the number of errors generated by an IP over a period of 10 se
27 stick-table type ip size 1m expire 10s store gpc0,http_err_rate(10s)
28 tcp-request connection track-sc1 src
29 tcp-request connection reject if { src_get_gpc0 gt 0 }
30
31 # Split static and dynamic traffic since these requests have different im
32 use_backend bk_web_static if { path_end .jpg .png .gif .css .js }
33
34 default_backend bk_web
35
36 # Dynamic part of the application
37 backend bk_web
38 balance roundrobin
39 cookie MYSRV insert indirect nocache
40
41 # If the source IP generated 10 or more http request over the defined per
42 # flag the IP as abuser on the frontend
43 acl abuse src_http_err_rate(ft_web) ge 10
44 acl flag_abuser src_inc_gpc0(ft_web)
45 tcp-request content reject if abuse flag_abuser
46
47 server srv1 192.168.1.2:80 check cookie srv1 maxconn 100
48 server srv2 192.168.1.3:80 check cookie srv2 maxconn 100
49
50 # Static objects
51 backend bk_web_static
52 balance roundrobin
53 server srv1 192.168.1.2:80 check maxconn 1000
54 server srv2 192.168.1.3:80 check maxconn 1000
```

```
ab -n 10 -c 10 http://127.0.0.1:8080/diskjirkuSjirkuSj
```

Watch the table content on the haproxy stats socket:

```
echo "show table ft_web" | socat unix:./haproxy.stats -  
# table: ft_web, type: ip, size:1048576, used:1  
0x8a9770: key=127.0.0.1 use=0 exp=5866 gpc0=1 http_err_rate(10000)=11
```

Let's try to run the same ab command and let's get the error:

```
apr_socket_recv: Connection reset by peer (104)
```

which means that HAProxy has blocked the IP address

## Notes

- We could combine configuration example above together to improve protection. This will be described later in an other article
- The numbers provided in the examples may be different for your application and architecture. Bench your configuration properly before applying in production.

## Related articles

- [Fight spam with early talking detection](#)
- [Protect Apache against Apache-killer script](#)
- [Protect your web server against slowloris](#)

## Links

- [HAProxy Technologies](#)
- [Aloha load balancer: HAProxy based LB appliance](#)

---

## 41 THOUGHTS ON “USE A LOAD-BALANCER AS A FIRST ROW OF DEFENSE AGAINST DDOS”



**dirigeant**

FEBRUARY 27, 2012 AT 16:45

Does “timeout http-request” affect POST requestst? 5 second may be too short for post requests. If it affects, is there any way to set timeout for only GET requests?



**Baptiste Assmann**

FEBRUARY 27, 2012 AT 16:53

Hi,

“timeout http-request” only applies to the header part of the request, and not to any data. As soon as the empty line is received, this timeout is not used anymore.

Cheers



**DmZ**

DECEMBER 12, 2012 AT 20:44

And how I could force timeout on data part of request? My goal to timeout slow http post attack (<http://www.darkreading.com/vulnerability-management/167901026/security/attacks-breaches/228000532/index.html>)



**Baptiste Assmann**

DECEMBER 13, 2012 AT 09:26

Hi DmZ,

Latest HAProxy version (1.5-dev15) allows layer 7 tracking so there may be some things to do with it to protect against this type of attack.

That said I'm not sure HAProxy has yet all the features required for this type of protection.

**dud**

OCTOBER 23, 2015 AT 19:08

Hi.

Looks like it is no longer true today as per your article :

<http://blog.haproxy.com/2015/10/14/whats-new-in-haproxy-1-6/>

“Once enabled, the timeout http-request parameters also apply to the POSTED data.”

BTW I did not find anything regarding this change in the latest changelog nor the latest documentation does point this out :

“Note that this timeout only applies to the header part of the request, and not to any data.”

Could you please clarify this change and possibly point me to the relevant message in the changelog ?

Thanks.

**Baptiste Assmann**

OCTOBER 27, 2015 AT 23:46

I confirm the documentation has not been updated accordingly. It will be updated to match real HAProxy's behavior.

Thanks for reporting.

**Patrick Mézard**

FEBRUARY 27, 2012 AT 22:51

“NOTE: if several users are hidden behind the same IP (NAT or proxy), this configuration may have a negative impact for them. You can whitelist these IPs.”

Not sure this is really practical for public web sites or mobile services. Or you have someone dedicated to whitelisting. He better be fast.

---

How do i get around this error

---



**Baptiste Assmann**

FEBRUARY 29, 2012 AT 10:22

Hi mandm,

You have to setup properly your stats socket in haproxy and point your socat to the socket path.

In the examples, the config file and the socket were in the same dir, which is not recommended in production. We usually configure the stat socket in /var/run.

Cheers

---



**Ivan Skyz**

FEBRUARY 29, 2012 AT 22:32

This is awesome, but is it possible to combine them all into one beautiful config under the same backend? I think I may speak for others when I find the syntax around the stick-table counters (gpc0) somewhat confusing.

---



**Baptiste Assmann**

MARCH 1, 2012 AT 07:27

Hi Ivan,

Yes it is possible.

I'll write this kind of conf in an other article, a bit later.

You can subscribe to the RSS stream or our tweeter account to get updated.

cheers

---

---

 **protection**

MARCH 8, 2012 AT 23:17

Thanks very interesting blog!

---

**Thomas**

JUNE 13, 2012 AT 12:03

Running 1.5dev11p20120604 i need to specify the sticktable in the frontend to make the reject work using the example of “Limiting the HTTP request rate”, instead of line 29:

```
tcp-request connection reject if { src_get_gpc0(ft_web) gt 0 }
```

---

**Vido**

JUNE 27, 2012 AT 15:18

With which version of haproxy is this possible? latest 1.5, or is it possible in version 1.4 as well? I kinda hate using development versions on production servers.

Thanks

---

**Baptiste Assmann**

JUNE 27, 2012 AT 22:01

Hi,

All the examples are related to 1.5 (dev) branch.

You're right, there are some 1.5 versions you should not use, like 1.5-dev9 and 1.5-dev10 🙄

To be honest, 1.5-dev7 is very stable and we 1.5-dev11 looks quite stable but is still young.

cheers

---

**Vido**

JUNE 28, 2012 AT 10:46

I see... When can we expect stable version?



---

keepalive on the server side, which requires huge modification on HAProxy's core.  
You can use 1.5-dev7 which is quite stable, I heard that the latest one, 1.5-dev11 is good as well.

cheers

---

Pingback: HTTP request flood mitigation | Exceliance – Aloha Load Balancer

---

Pingback: Scalable WAF protection with HAProxy and Apache with modsecurity | Exceliance – Aloha Load Balancer

---

Pingback: high performance WAF platform with Naxsi and HAProxy | Exceliance – Aloha Load Balancer



**Smana**

MAY 2, 2013 AT 18:03

Hello, first of all Good job 😊 !

I just wanted to know how is it possible to reduce the number of lines if we want to use each of the configuration you proposed.

For example is it possible to have one line for the sticky-table like the following one ? (i know that one doesn't seem to work :p )

```
stick-table type ip size 1m expire 30s store gpc0,http_req_rate(10s),http_err_rate(10s) store  
conn_cur store conn_rate(3s)
```

Thanks,  
Smana



**Baptiste Assmann**



**LRaikhman (@lraikhman)**

MAY 13, 2013 AT 19:05

Hi,

I want to implement the vulnerability scan detection (your last example) but want to exclude one IP address from the detection.

Can you help me to do that – is that possible?



**Baptiste Assmann**

MAY 14, 2013 AT 11:22

Hi,

Yes it is possible.

Look for the whitelisting options proposed in some of the configuration example.

Baptiste



**Rfraile**

JULY 17, 2013 AT 11:48

Hello Baptiste,

Why you define a “tcp\_max\_syn\_backlog” with syncookies enabled?

In this situation, the backlog is not used because the aren’t any entry in that table, isn’t it?

Thanks,



**Rfraile**

JULY 17, 2013 AT 12:30

**Nathan**

MARCH 3, 2014 AT 11:16

Hello,

Thanks for your topic.

I have a warning in HAPROXY 1.5.-dev21 :

Starting HAproxy: [WARNING] 061/100656 (6315) : parsing acl keyword  
'src\_inc\_gpc0(HTTP\_FR\_PHP55)' :  
no pattern to match against were provided, so this ACL will never match.  
If this is what you intended, please add '-' to get rid of this warning.  
If you intended to match only for existence, please use '-m found'.  
If you wanted to force an int to match as a bool, please use '-m bool'.

I don't understand because on a 1.5-dev19, i don't have this warning.

Thanks for your help

**Baptiste Assmann**

MARCH 15, 2014 AT 22:13

Hi,

Please send your question to the ML, including your configuration.

Baptiste

---

Pingback: Scalable WAF protection with HAProxy and Apache with modsecurity | HAProxy Technologies – Aloha Load Balancer

---

Pingback: high performance WAF platform with Naxsi and HAProxy | HAProxy Technologies – Aloha Load Balancer

---

For keep-alive clients it is also convenient to rate limit session rate.  
Very nice writeup, thank you!

---

**om**

MAY 21, 2014 AT 07:51

Hi Can it be possible to throttle limit the HTTP Post ad Get method at HAProxy layer ?

---

**Baptiste Assmann**

MAY 21, 2014 AT 08:17

Yes of course.

With HAProxy, you have ACLs to match HTTP methods.

Baptiste

---

**om**

MAY 21, 2014 AT 08:56

Sorry, i am not getting the complete command. could you please provide me. i found following sample which block the http request if it not belongs to GET/POST/OPTIONS method:

```
acl missing_cl hdr_cnt(Content-length) eq 0
block if HTTP_URL_STAR !METH_OPTIONS || METH_POST missing_cl
block if METH_GET HTTP_CONTENT
block unless METH_GET or METH_POST or METH_OPTIONS
```

Best Regards

-Om

---

**Tunisie annonces**

NOVEMBER 28, 2015 AT 19:56

**Baptiste Assmann**

NOVEMBER 30, 2015 AT 07:54

Hi,

Yes, you can do this, using `req.hdr(Referrer)`.

**Alexey**

APRIL 28, 2016 AT 09:43

Hi folks,

As far as i understand, `src_http_err_rate` does NOT counts 5xx http errors, right?

Currently I am under DOS attack which causes a lot of 500 errors from the web application, and error counter is incrementing only on 4xx.

Is there any option how can i track 5xx errors and do “tcp-request reject ” if greater than some value?

Thanks in advance!

**Baptiste Assmann**

MAY 2, 2016 AT 07:52

You're right. This counter does not count HTTP 500 errors.

You may want to use `gpc0`, increase `gpc0` on responses where status is greater or equal than 500.

and then decide to deny if the `gpc0` inc rate is greater than some threshold.

**Sean**

AUGUST 4, 2016 AT 15:43

I have an HAProxy for a mail setup. Unfortunately sometimes accounts get hacked and we get a lot of spam and smtp connections coming through.

Usually, they are all from the same IP address. Would this config help me control it?

---

```
tcp-request connection reject if { sc1_conn_rate ge 20 }
tcp-request connection reject if { sc1_conn_cur ge 20 }
tcp-request connection track-sc1 src
acl local_ips src -f /etc/haproxy/trusted-ips.txt
use_backend smtp-be-local if local_ips
default_backend smtp-be-foreign
```

```
backend smtp-be-foreign
option smtpchk
source 0.0.0.0 usesrc clientip
server xxxx01 x.x.x.x:25 maxconn 1000 check port 25
```

```
backend smtp-be-local
option smtpchk
source 0.0.0.0 usesrc clientip
server xxxx02 x.x.x.x:25 maxconn 1000 check port 25
server xxxx03 x.x.x.x:25 maxconn 1000 check port 25
```

Any other suggestions?